

Data Warehouse Assignment

Lab 4

Gyanendra Kr Shukla
191112040
CSE-1

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

titanic_df = pd.read_excel('titanic.xls')
print("Describing the dataframe:")
titanic_df.describe()
```

Describing the dataframe:

	pclass	survived	age	sibsp	parch \
count	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000
mean	2.294882	0.381971	29.881135	0.498854	0.385027
std	0.837836	0.486055	14.413500	1.041658	0.865560
min	1.000000	0.000000	0.166700	0.000000	0.000000
25%	2.000000	0.000000	21.000000	0.000000	0.000000
50%	3.000000	0.000000	28.000000	0.000000	0.000000
75%	3.000000	1.000000	39.000000	1.000000	0.000000
max	3.000000	1.000000	80.000000	8.000000	9.000000

	fare
count	1308.000000
mean	33.295479
std	51.758668
min	0.000000
25%	7.895800
50%	14.454200
75%	31.275000
max	512.329200

```
print("First 5 rows of the dataframe:")
titanic_df.head()
```

First 5 rows of the dataframe:

	pclass	survived	name	sex \
0	1	1	Allen, Miss. Elisabeth Walton	female
1	1	1	Allison, Master. Hudson Trevor	male
2	1	0	Allison, Miss. Helen Loraine	female
3	1	0	Allison, Mr. Hudson Joshua Creighton	male
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female

	age	sibsp	parch	ticket	fare	embarked
0	29.0000	0	0	24160	211.3375	S
1	0.9167	1	2	113781	151.5500	S
2	2.0000	1	2	113781	151.5500	S
3	30.0000	1	2	113781	151.5500	S
4	25.0000	1	2	113781	151.5500	S

Q1. WAP to compute following on the titanic dataset

1. Counting missing values in each column
2. Removing rows/attributes with missing values and calculating the ratio with respect to original data
3. Replacing missing values of the age attribute with
 - (a) Mean
 - (b) Median

```
# Counting missing values in each column
nas = {}
for series in titanic_df:
    na = 0
    for value in titanic_df[series]:
        # print(value, end=', ')
        if pd.isnull(value):
            na += 1
    nas[series] = na
null_vals = pd.Series(nas, name="No. of null values")

# Removing rows with missing values
titanic_row_null_removed = titanic_df.copy()
null = []
for row in titanic_df.itertuples():
    for val in list(row)[1:]:
        if pd.isnull(val):
            null.append(row[0])
            break

for row in null:
    titanic_row_null_removed.drop(row, inplace=True)
print(f"Ratio of original with null removed dataframe: {titanic_row_null_removed.shape[0]/titanic_df.shape[0]}")
Ratio of original with null removed dataframe: 0.7967914438502673

# removing columns with missing values
titanic_col_null_removed = titanic_df.copy()
null = []
for col in titanic_df.columns:
```

```

    for val in titanic_df[col].tolist():
        if pd.isnull(val):
            null.append(col)
            break
for col in null:
    titanic_col_null_removed.drop(col, axis=1, inplace=True)
print(f"Ratio of original with null removed dataframe: {titanic_col_null_removed.shape[1]/titanic_df.shape[1]}")
Ratio of original with null removed dataframe: 0.7

# replacing missing values of attribute with mean
titanic_df['meanfilled-age'] = titanic_df['age'].fillna(titanic_df['age'].mean())

```

```
titanic_df.describe()
```

	pclass	survived	age	sibsp	parch \
count	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000
mean	2.294882	0.381971	29.881135	0.498854	0.385027
std	0.837836	0.486055	14.413500	1.041658	0.865560
min	1.000000	0.000000	0.166700	0.000000	0.000000
25%	2.000000	0.000000	21.000000	0.000000	0.000000
50%	3.000000	0.000000	28.000000	0.000000	0.000000
75%	3.000000	1.000000	39.000000	1.000000	0.000000
max	3.000000	1.000000	80.000000	8.000000	9.000000

	fare	meanfilled-age
count	1308.000000	1309.000000
mean	33.295479	29.881135
std	51.758668	12.883199
min	0.000000	0.166700
25%	7.895800	22.000000
50%	14.454200	29.881135
75%	31.275000	35.000000
max	512.329200	80.000000

```

# replacing missing values of attribute with median
titanic_df['medianfilled-age'] = titanic_df['age'].fillna(titanic_df['age'].median())

```

```
titanic_df.describe()
```

	pclass	survived	age	sibsp	parch \
count	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000
mean	2.294882	0.381971	29.881135	0.498854	0.385027
std	0.837836	0.486055	14.413500	1.041658	0.865560
min	1.000000	0.000000	0.166700	0.000000	0.000000
25%	2.000000	0.000000	21.000000	0.000000	0.000000
50%	3.000000	0.000000	28.000000	0.000000	0.000000
75%	3.000000	1.000000	39.000000	1.000000	0.000000

max	3.000000	1.000000	80.000000	8.000000	9.000000
-----	----------	----------	-----------	----------	----------

	fare	meanfilled-age	medianfilled-age
count	1308.000000	1309.000000	1309.000000
mean	33.295479	29.881135	29.503183
std	51.758668	12.883199	12.905246
min	0.000000	0.166700	0.166700
25%	7.895800	22.000000	22.000000
50%	14.454200	29.881135	28.000000
75%	31.275000	35.000000	35.000000
max	512.329200	80.000000	80.000000

Q2. WAP to perform transformation of data on age attribute of Titanic dataset using Binning

1. With equal width bins using K=25, using mean (you may fill zero value to balance last bin)
2. With equal frequency bins of size 100 using mean (you may fill zero value to balance last bin)

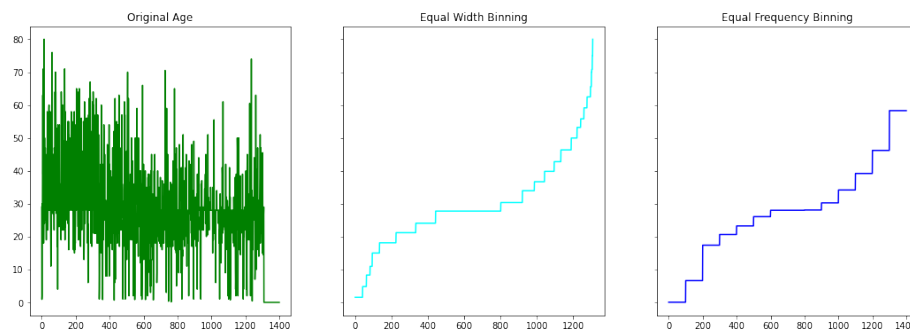
```
def binning_equal_width(ageser, k):
    w = (max(ageser) - min(ageser))/k
    bins = [min(ageser) + i*w for i in range(k+1)]
    ret = []
    for i in range(0, k):
        curr = []
        for j in ageser:
            if j >= bins[i] and j <= bins[i+1]:
                curr.append(j)
        # mean of curr
        ret+=[np.mean(curr)]*len(curr)
    return ret

def binning_equal_frequency(age, size):
    while len(age) % size != 0:
        age.append(0)
    age = sorted(age)
    ret = []
    for i in range(0, len(age), size):
        ret+=[np.mean(age[i:i+size])] * size
    return ret

def binning(series):
    equalWidth = binning_equal_width(series, 25)
    equalFrequency = binning_equal_frequency(series, 100)
```

```
_, axs = plt.subplots(1,3, figsize=(18,6), sharey=True)
axs[0].plot(series, label = 'Original Age', color='green')
axs[0].set_title('Original Age')
axs[1].plot(equalWidth, label = 'Equal Width Binning', color = 'cyan')
axs[1].set_title('Equal Width Binning')
axs[2].plot(equalFrequency, label = 'Equal Frequency Binning', color = 'blue')
axs[2].set_title('Equal Frequency Binning')
plt.show()
```

```
binning(titanic_df['medianfilled-age'].tolist())
```



Q3. Performing Chi-Square test on the titanic dataset (and show the contingency table)

1. Pclass & Survived (fill missing values with highest frequency category, if any)
2. Pclass & sex (fill missing values with highest frequency category, if any)
3. Pclass & embarked (fill missing values with highest frequency category, if any)

```
import seaborn as sns
from scipy.stats.distributions import chi2
def chiSquareCalc(x, y):
    x = x.fillna(x.mode()[0])
    y = y.fillna(y.mode()[0])

    table = {}
    for i in range(len(x)):
        if x[i] in table:
            if y[i] in table[x[i]]:
                table[x[i]][y[i]] += 1
            else:
                table[x[i]][y[i]] = 1
        else:
            table[x[i]] = {}
```

```

        table[x[i]][y[i]] = 1
df = pd.DataFrame(table)
df['Sum(row)'] = df.sum(axis=1)
df.loc['Sum(col)'] = df.sum()
df.columns.name = x.name
df.index.name = y.name

df_exp = df.copy()


for i in range(len(df.columns)-1):
    df_exp.iloc[i, :-1] = df_exp.iloc[-1, :-1]*df_exp.iloc[i, -1]/df_exp.iloc[-1, -1]

# Plot everything
plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
sns.heatmap(df.iloc[:-1, :-1], annot=True, fmt='d')
plt.title('Contingency Table')
plt.subplot(1,2,2)
sns.heatmap(df_exp.iloc[:-1, :-1], annot=True, fmt='g')
plt.title('Expected Value Table')
plt.tight_layout()
plt.show()

chiSq = np.sum(np.sum(((df.iloc[:-1, :-1] - df_exp.iloc[:-1, :-1])**2/df_exp.iloc[:-1, :-1],
print("The chi-square value:", chiSq)
rows = x.unique()
cols = y.unique()
dof = (len(rows)-1)*(len(cols)-1)
print("The degree of freedom:", dof)
p = chi2.sf(chiSq, dof)
print("The p-value:", p)

chiSquareCalc(titanic_df['pclass'], titanic_df['survived'])

```

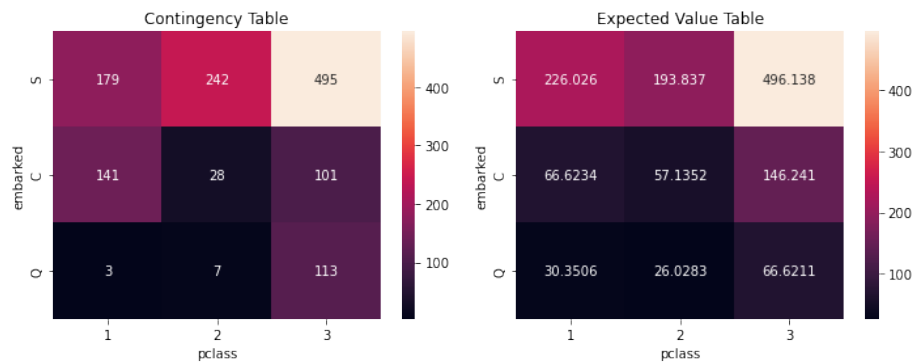


The chi-square value: 127.85915643930326

The degree of freedom: 2

The p-value: 1.7208259588256175e-28

```
chiSquareCalc(titanic_df['pclass'], titanic_df['embarked'])
```




The chi-square value: 204.48431967559742

The degree of freedom: 4

The p-value: 4.0799162291284984e-43

```
chiSquareCalc(titanic_df['pclass'], titanic_df['sex'])
```



The chi-square value: 20.378781205085584

The degree of freedom: 2

The p-value: 3.7566772719164106e-05

Q4. Perform correlation analysis

1. Age & Fare (first fill missing values using mean, if any)
2. Age & sibsp (first fill missing values using mean, if any)


```
def correlation(x, y):  
    x = x.fillna(x.mean())  
    y = y.fillna(y.mean())
```

```

r = ((x-x.mean()*(y-y.mean())).sum() / ((x.size-1)*x.std()*y.std()))
plt.scatter(x, y, c=np.random.rand(x.size))
plt.title("Correlation: "+str(r), fontsize=15)
plt.show()

correlation(titanic_df['age'], titanic_df['fare'])

```



```

correlation(titanic_df['age'], titanic_df['sibsp'])

```